

In [1]:

```
# In[1]:

#Import environments:

import pandas as pd
import numpy as np
import sklearn
import os
from sklearn import linear_model
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, roc_curve
import seaborn as sns
import matplotlib.pyplot as plt
```

In [6]:

```
# In[ ]:

#Load Data:

airport = pd.read_excel(os.path.expanduser("/Users/carlosalvarez/Desktop/12Months.xl
```

In [7]:

```
# In[ ]:

airport.shape
```

Out[7]:

(100000, 20)

In [8]:

```
# In[ ]:

airport.isnull().values.any()
```

Out[8]:

True

In [9]:

```
# In[ ]:
```

```
airport.isnull().sum()
```

Out[9]:

```
YEAR                0
QUARTER             0
MONTH              0
DAY_OF_MONTH        0
ORIGIN              0
ORIGIN_CITY_NAME    0
ORIGIN_STATE_ABR    0
DEST               0
DEST_CITY_NAME      0
DEST_STATE_ABR      0
CRS_DEP_TIME        0
DEP_TIME            1258
DEP_DELAY           1511
DEP_DEL15           1511
CRS_ARR_TIME        0
ARR_TIME            1350
ARR_DELAY           1639
ARR_DEL15           1639
CANCELLED           0
DISTANCE            0
dtype: int64
```

In [10]:

```
# In[ ]:
```

```
columns_to_drop = ["ORIGIN_STATE_ABR", "DEST_CITY_NAME", "DEST_STATE_ABR", "ORIGIN_CITY_NAME"]
```

In [11]:

```
# In[ ]:
```

```
airport.fillna(value=0, inplace=True)
airport.isnull().values.any()
```

Out[11]:

False

In []:

In [12]:

```
# In[ ]:

airport = pd.get_dummies(airport, columns=[ 'ORIGIN' , 'DEST' ])
airport.head()
```

Out[12]:

	YEAR	QUARTER	MONTH	DAY_OF_MONTH	ORIGIN_CITY_NAME	ORIGIN_STATE_ABR	DEST_
0	2018	1	1	27	Fort Lauderdale, FL	FL	
1	2018	1	1	27	Seattle, WA	WA	San F
2	2018	1	1	27	Washington, DC	VA	
3	2018	1	1	27	Los Angeles, CA	CA	
4	2018	1	1	27	Jacksonville, FL	FL	

5 rows × 403 columns

In [13]:

```
# In[ ]:

airport.drop(labels=columns_to_drop, axis=1, inplace=True)
train_x, test_x, train_y, test_y = train_test_split(airport.drop('ARR_DEL15', axis=1),
```

In [14]:

```
# In[ ]:

train_x.shape
```

Out[14]:

(80000, 398)

In [15]:

```
# In[ ]:
```

```
test_x.shape
```

Out[15]:

```
(20000, 398)
```

In [271]:

```
# In[ ]:
```

```
gboost = GradientBoostingClassifier(criterion='friedman_mse', init=None,
    learning_rate=0.085, loss='deviance', max_depth=5,
    max_features=8, max_leaf_nodes=None,
    min_impurity_decrease=0.0, min_impurity_split=None,
    min_samples_leaf=1, min_samples_split=2,
    min_weight_fraction_leaf=0.0, n_estimators=100,
    n_iter_no_change=None, presort='auto', random_state=None,
    subsample=1.0, tol=0.0001, validation_fraction=0.1,
    verbose=0, warm_start=False)
gboost.fit(train_x, train_y)
```

Out[271]:

```
GradientBoostingClassifier(criterion='friedman_mse', init=None,
    learning_rate=0.085, loss='deviance', max_depth=5,
    max_features=8, max_leaf_nodes=None,
    min_impurity_decrease=0.0, min_impurity_split=None,
    min_samples_leaf=1, min_samples_split=2,
    min_weight_fraction_leaf=0.0, n_estimators=100,
    n_iter_no_change=None, presort='auto', random_state=None
,
    subsample=1.0, tol=0.0001, validation_fraction=0.1,
    verbose=0, warm_start=False)
```

In [272]:

```
# In[ ]:
```

```
y_gboost_pred = gboost.predict(test_x)
```

```
labels = [0, 1]
```

```
cm = confusion_matrix(test_y, y_gboost_pred, labels)
```

```
gboost_accuracy = str(np.round(100*float(cm[0][0]+cm[1][1])/float((cm[0][0]+cm[1][1]
gboost_recall = str(np.round(100*float((cm[1][1]))/float((cm[1][0]+cm[1][1])),2))
```

```
print('Accuracy: ' + gboost_accuracy + '%')
```

```
print('Recall: ' + gboost_recall + '%')
```

```

print('Confusion matrix:')
print(cm)


fpr, tpr, _ = roc_curve(test_y, y_gboost_pred)
auc = np.trapz(fpr,tpr)
print('Area under the ROC curve: ' + str(auc))

fig = plt.figure(1)
plt.plot(fpr,tpr,color='green')
plt.xlabel('False positive rate (FPR)')
plt.ylabel('True positive rate (TPR)')
plt.title('Receiver operating characteristic (ROC)')


fig = plt.figure(2)
ax = fig.add_subplot(111)
cax = ax.matshow(cm)
plt.title('Confusion matrix for Gradient Boosting classifier with original data')
fig.colorbar(cax)
ax.set_xticklabels([''] + labels)
ax.set_yticklabels([''] + labels)
plt.xlabel('Predicted')
plt.ylabel('True')
plt.show()


```

Accuracy: 96.88%
 Recall: 84.63%
 Confusion matrix:
 [[15951 2]
 [622 3425]]
 Area under the ROC curve: 0.07690973133009896



Confusion matrix for Gradient Boosting classifier with original data





In [269]:

```
# In[ ]:
```

```
dtree = DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=15,
                               max_features=5, max_leaf_nodes=None,
                               min_impurity_decrease=0.0, min_impurity_split=None,
                               min_samples_leaf=1, min_samples_split=2,
                               min_weight_fraction_leaf=0.0, presort=False, random_state=None,
                               splitter='best')
dtree.fit(train_x, train_y)
```

Out[269]:

```
DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=
15,
                       max_features=5, max_leaf_nodes=None, min_impurity_decrease
=0.0,
                       min_impurity_split=None, min_samples_leaf=1,
                       min_samples_split=2, min_weight_fraction_leaf=0.0,
                       presort=False, random_state=None, splitter='best')
```

In [270]:

```
# In[ ]:
```

```
y_dtree_pred = dtree.predict(test_x)

labels = [0, 1]
cm = confusion_matrix(test_y, y_dtree_pred, labels)

dtree_accuracy = str(np.round(100*float(cm[0][0]+cm[1][1])/float((cm[0][0]+cm[1][1]
dtree_recall = str(np.round(100*float((cm[1][1]))/float((cm[1][0]+cm[1][1])),2))

print('DT Accuracy: ' + dtree_accuracy + '%')
print('Recall: ' + dtree_recall + '%')
print('Confusion matrix:')
print(cm)

fpr, tpr, _ = roc_curve(test_y, y_dtree_pred)
auc = np.trapz(fpr,tpr)
print('Area under the ROC curve: ' + str(auc))
```


```

fig = plt.figure(1)
plt.plot(fpr,tpr,color='green')
plt.xlabel('False positive rate (FPR)')
plt.ylabel('True positive rate (TPR)')
plt.title('Receiver operating characteristic (ROC)')

fig = plt.figure(2)
ax = fig.add_subplot(111)
cax = ax.matshow(cm)
plt.title('Confusion matrix for Decision Tree classifier with original data')
fig.colorbar(cax)
ax.set_xticklabels([''] + labels)
ax.set_yticklabels([''] + labels)
plt.xlabel('Predicted')
plt.ylabel('True')
plt.show()


```

DT Accuracy: 97.6%
Recall: 95.58%
Confusion matrix:
[[15652 301]
[179 3868]]
Area under the ROC curve: 0.031549109286636734



Confusion matrix for Decision Tree classifier with original data





In [228]:

```
# In[ ]:
```

```
rforest = RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                                max_depth=6, max_features='auto', max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=4,
                                min_weight_fraction_leaf=0.0, n_estimators=5, n_jobs=None,
                                oob_score=False, random_state=None, verbose=0,
                                warm_start=False)
rforest.fit(train_x, train_y)
```

Out[228]:

```
RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                        max_depth=6, max_features='auto', max_leaf_nodes=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=4,
                        min_weight_fraction_leaf=0.0, n_estimators=5, n_jobs=None,
                        oob_score=False, random_state=None, verbose=0,
                        warm_start=False)
```

In [232]:

```
# In[ ]:
```

```
y_rforest_pred = rforest.predict(test_x)

labels = [0, 1]
cm = confusion_matrix(test_y, y_rforest_pred, labels)

rforest_accuracy = str(np.round(100*float(cm[0][0]+cm[1][1])/float((cm[0][0]+cm[1][1]+cm[0][1]+cm[1][0])),2))
rforest_recall = str(np.round(100*float((cm[1][1]))/float((cm[1][0]+cm[1][1])),2))

print('RF Accuracy: ' + rforest_accuracy + '%')
print('Recall: ' + rforest_recall + '%')
print('Confusion matrix:')
print(cm)

fpr, tpr, _ = roc_curve(test_y, y_rforest_pred)
auc = np.trapz(fpr,tpr)
print('Area under the ROC curve: ' + str(auc))
```

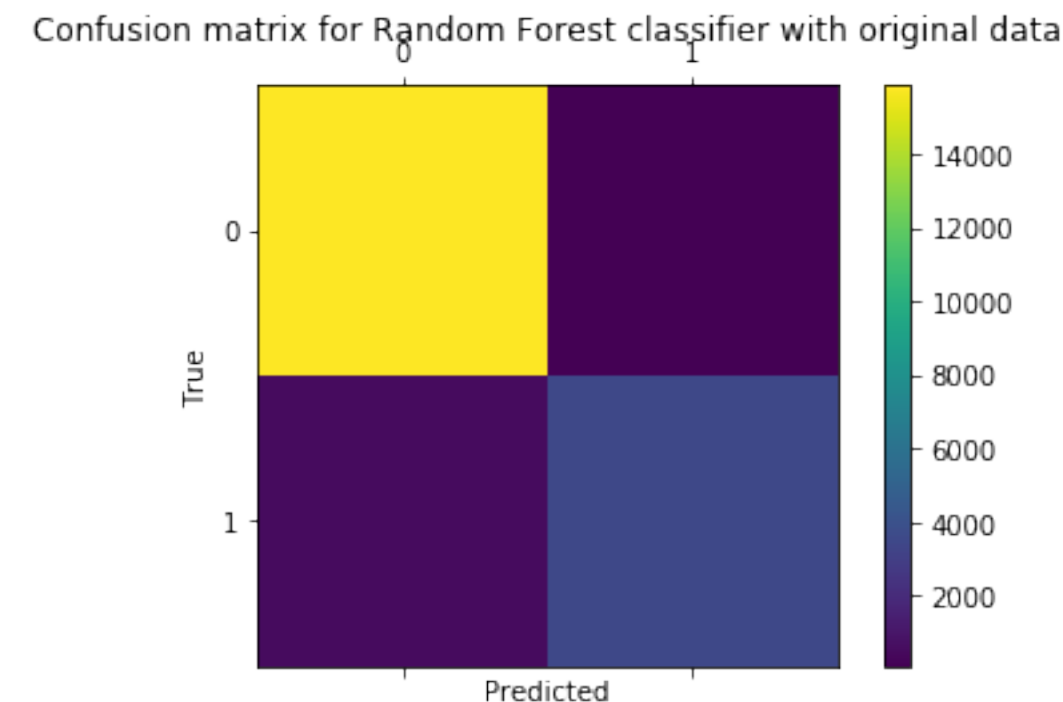

```
fig = plt.figure(1)
plt.plot(fpr,tpr,color='green')
plt.xlabel('False positive rate (FPR)')
```



```
plt.xlabel('False positive rate (FPR)')
plt.ylabel('True positive rate (TPR)')
plt.title('Receiver operating characteristic (ROC)')

fig = plt.figure(2)
ax = fig.add_subplot(111)
cax = ax.matshow(cm)
plt.title('Confusion matrix for Random Forest classifier with original data')
fig.colorbar(cax)
ax.set_xticklabels([''] + labels)
ax.set_yticklabels([''] + labels)
plt.xlabel('Predicted')
plt.ylabel('True')
plt.show()
```

RF Accuracy: 96.91%
Recall: 86.76%
Confusion matrix:
[[15871 82]
[536 3511]]
Area under the ROC curve: 0.06879194228053555



In [22]:

```
reg = linear_model.LogisticRegression(C=1.0, class_weight=None, dual=True, fit_intercept=True,
    intercept_scaling=10, max_iter=20, multi_class='warn',
    n_jobs=None, penalty='l2', random_state=2, solver='warn',
    tol=0.008, verbose=90, warm_start=True)
reg.fit(train_x, train_y)
```

```
/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.
  FutureWarning)
```

[LibLinear]

```
/anaconda3/lib/python3.7/site-packages/sklearn/svm/base.py:922: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.
  "the number of iterations.", ConvergenceWarning)
```

Out[22]:

```
LogisticRegression(C=1.0, class_weight=None, dual=True, fit_intercept=True,
    intercept_scaling=10, max_iter=20, multi_class='warn',
    n_jobs=None, penalty='l2', random_state=2, solver='warn',
    tol=0.008, verbose=90, warm_start=True)
```

In [23]:

```
reg = linear_model.LogisticRegression(C=1.0, class_weight=None, dual=True, fit_intercept=True,
    intercept_scaling=10, max_iter=20, multi_class='warn',
    n_jobs=None, penalty='l2', random_state=2, solver='warn',
    tol=0.008, verbose=90, warm_start=True)
reg.fit(train_x, train_y)

y_reg_pred = reg.predict(test_x)

labels = [0, 1]
cm = confusion_matrix(test_y, y_reg_pred, labels)

reg_accuracy = str(np.round(100*float(cm[0][0]+cm[1][1])/float((cm[0][0]+cm[1][1] + cm[0][1]+cm[1][0]),2))
reg_recall = str(np.round(100*float((cm[1][1]))/float((cm[1][0]+cm[1][1])),2))

print('Accuracy: ' + reg_accuracy + '%')
print('Recall: ' + reg_recall + '%')
print('Confusion matrix:')
print(cm)

fpr, tpr, _ = roc_curve(test_y, y_reg_pred)
auc = np.trapz(fpr,tpr)
print('Area under the ROC curve: ' + str(auc))
```

```
fig = plt.figure(1)
plt.plot(fpr,tpr,color='green')
plt.xlabel('False positive rate (FPR)')
plt.ylabel('True positive rate (TPR)')
plt.title('Receiver operating characteristic (ROC)')
```

```
fig = plt.figure(2)
ax = fig.add_subplot(111)
cax = ax.matshow(cm)
plt.title('Confusion Matrix for Linear Regression')
fig.colorbar(cax)
ax.set_xticklabels([''] + labels)
ax.set_yticklabels([''] + labels)
plt.xlabel('Predicted')
plt.ylabel('True')
plt.show()
```

```
/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/logistic.p
y:433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.2
2. Specify a solver to silence this warning.
```

```
FutureWarning)
```

```
[LibLinear]
```

```
/anaconda3/lib/python3.7/site-packages/sklearn/svm/base.py:922: Conver
genceWarning: Liblinear failed to converge, increase the number of ite
rations.
```

```
"the number of iterations.", ConvergenceWarning)
```


```
Accuracy: 98.5%
```

```
Recall: 92.61%
```

```
Confusion matrix:
```

```
[[15951    2]
 [ 299  3748]]
```


```
Area under the ROC curve: 0.03700362804371396
```



Confusion Matrix for Linear Regression

0

1



In []:

In []:

In []:

In [1]:

In [6]:

In [7]:

Out[7]:
(100000, 20)

In [8]:

Out[8]:
True